

# CMSC 424 First Midterm

Huang Yufan

TOTAL POINTS

**35.5 / 50**

QUESTION 1

**1 Q1 1 / 1**

- **1 pts** Incorrect (Should be DDL)
- **0.75 pts** Left blank (get 25%)
- ✓ - **0 pts** Correct

QUESTION 2

**2 Q2 2 / 2**

- ✓ - **0 pts** Correct
- **1.5 pts** No answer
- **1 pts** Incorrect query syntax
- **0.5 pts** password is a string
- **0.5 pts** Double quotes or no quotes
- **0.25 pts** Double quotes or no quotes
- **0.5 pts** user\_id is an int
- **2 pts** incorrect

QUESTION 3

**3 Q3 2 / 3**

- **0 pts** Correct
- ✓ - **1 pts** Invalid Syntax
- **0.5 pts** Minor Issue
- **1 pts** Not necessarily unique
- **1.5 pts** Didn't Use Update
- **2.25 pts** Blank / Minor partial credit
- **3 pts** No partial credit

QUESTION 4

**4 Q4 2 / 2**

- ✓ - **0 pts** Correct
- **1 pts** Where clause is incorrect
- **1 pts** Incorrect delete clause
- **1.5 pts** 1/4 of the points for not answering
- **0.5 pts** Minor syntax issue (example: using double quotes instead of single or including \* in the delete

clause)

- **0.5 pts** Second minor syntax issue

QUESTION 5

**5 Q5a 1 / 1**

- ✓ - **0 pts** Correct
- **1 pts** Incorrect

QUESTION 6

**6 Q5b 1 / 1**

- ✓ - **0 pts** Correct
- **1 pts** Incorrect
- **0.75 pts** Left blank (get 25%)

QUESTION 7

**7 Q5c 1 / 1**

- ✓ - **0 pts** Correct
- **1 pts** Incorrect
- **0.75 pts** Left Blank (get 25%)

QUESTION 8

**8 Q6 1.5 / 3**

- + **3 pts** Correct
- + **2.5 pts** Using like on wrong attribute (user\_id)
- ✓ + **1.5 pts** Using = instead of like
- + **1.5 pts** Using in matches any email that contains 'edu' in any position.
- + **2.5 pts** Minor Error :Wrong usage of like. Used 'edu%' or '[\s\d]\*edu' instead of '%edu'
- + **3 pts** Return only user\_id. Not all attributes
- + **1 pts** Major Syntax Issue
- + **0.75 pts** Completely Blank
- + **0 pts** No partial credit
- + **2 pts** Almost correct with minor error

QUESTION 9

## 9 Q7 2 / 4

- 0 pts Correct
- ✓ - 2 pts Should only counts notes, not comments (i.e., parent\_id is NULL)
- 2 pts Should not be a moderator (for any topic)
- 4 pts Incorrect (too far from correct solution to be able to assign partial credit)
- 1 pts SQL minor issue: attribute name for author in Posts should be author\_id; in Moderators it should be user\_id
- 1 pts SQL minor issue: ambiguous attribute names, user\_id.
- 1 pts SQL minor issue: '(Not) In' should be used to test set containment (i.e., p.author\_id NOT IN (select user\_id from moderators) )
- 1 pts SQL minor issue: with clause used as nested query
- 2 pts Used 'Or' instead of 'And': Should be notes and not by a moderator at the same time.

## QUESTION 10

### 10 Q8 3 / 4

- 0 pts Correct
- ✓ - 1 pts Almost (ans is BCD)
- 2 pts Missed two of them (correct ans is BCD)
- 3 pts Missed 3 out of 4 or left blank (correct ans is BCD)
- 4 pts Opposite of what is correct (BCD)

## QUESTION 11

### 11 Q9 1 / 4.5

- 0 pts Correct
- 0.5 pts Very, very close (correct ans is DF)
- 1 pts Almost (correct ans is DF)
- 1.5 pts Close (correct ans is DF)
- 2 pts Several mistakes (correct ans is DF)
- 2.5 pts More incorrect than correct (correct ans is DF)
- 3 pts Some correct (correct ans is DF)
- ✓ - 3.5 pts Mostly incorrect (correct ans is DF)
- 3.38 pts Left blank (get 25% credit)
- 4 pts Not close (correct ans is DF)

## QUESTION 12

### 12 Q10 3 / 4

- 0 pts Correct
- ✓ - 1 pts Partly correct -- correlated correctly on topic\_id, but not on author\_id.
- 1 pts Partly correct -- correlated correctly on author\_id, but not on topic\_id.
- 2 pts Partly correct but no correlated query.
- 2 pts Partly correct. Syntax errors and correlated query not used correctly.
- 2 pts Partly correct, one or both of author\_id and topic\_id not considered for correlation.
- 3 pts Incorrect, but thought along the right direction
- 3 pts No answer / left blank, get 25% credit
- 4 pts Incorrect (too far away to be able to justify partial credit)

## QUESTION 13

### 13 Q11 2 / 2

- 0 pts Correct
- ✓ - 0 pts Incorrect (technically possible if, e.g. (null, 1) is in moderators), but only taking off one point.

## QUESTION 14

### 14 Q12a 0 / 1

- 0 pts Correct
- ✓ - 1 pts Incorrect (should be 1)
- 0.75 pts Question asked how many rows, not value that is returned.

## QUESTION 15

### 15 Q12b 1 / 1

- ✓ - 0 pts Correct
- 1 pts Incorrect
- 0 pts Question asked how many rows, not what is returned, but will give full credit in this

## QUESTION 16

### 16 Q12c 0 / 1

- 0 pts Correct
- ✓ - 1 pts Incorrect (correct answer is 4)

- **0 pts** Question asked how many rows, not output, but will give full credit in this case  
- **0.75 pts** Blank

#### QUESTION 17

#### 17 Q12d 2 / 2

✓ - **0 pts** Correct  
- **2 pts** Incorrect (correct answer is 0)  
- **1.5 pts** Blank (get 25% credit)

#### QUESTION 18

#### 18 Q12e 1 / 1

✓ - **0 pts** Correct  
- **1 pts** Incorrect (correct answer is 2)  
- **0.75 pts** Blank (get 25% credit)  
- **0 pts** Q asked how many rows, but will mark correct in this case

#### QUESTION 19

#### 19 Q12f 0 / 1

- **0 pts** Correct  
✓ - **1 pts** Incorrect (correct answer is 3)  
- **0.75 pts** Blank (get 25% credit)  
- **0 pts** Q asked for num rows, not output, but will mark correct in this case

#### QUESTION 20

#### 20 Q13 4 / 4

✓ + **4 pts** Correct  
+ **3 pts** Partially correct. Need to add OR between INSERT and UPDATE  
+ **2 pts** Partially correct. Need to add "OR UPDATE"  
+ **0 pts** Incorrect / too far away to justify partial credit  
+ **1 pts** Need to add the correct event(s). i.e. INSERT / UPDATE  
+ **1 pts** Left blank, get 25% credit  
+ **3 pts** Syntax Issue. You need to use BEFORE once.  
+ **2 pts** Partially correct.  
+ **1 pts** Not quite there.  
+ **2 pts** Syntax Issue. You need to use BEFORE

once and use OR instead of ,

+ **2 pts** Syntax Issue. You need to use BEFORE once and use OR instead of AND

+ **3.25 pts** You need to add BEFORE instead of AFTER

+ **3 pts** Need UPDATE as well  
+ **2 pts** Partially correct.

#### QUESTION 21

#### 21 Q14 5 / 6.5

✓ + **1 pts** Part 1: single line between users and moderators  
✓ + **1 pts** Part 1: no arrows on line between users and moderators  
✓ + **1 pts** Part 1: double line between topics and moderators  
✓ + **1 pts** Part 1: no arrows on line between topics and moderators  
+ **1.5 pts** Part 2: no attributes added to the moderators relationship set  
✓ + **0.5 pts** Part 3: only user\_id underlined in users entity set  
✓ + **0.5 pts** Part 3: only topic\_id underlined in topics entity set

Name: yufan Huang

**UMD College Park Department of Computer-Science  
CMSC 424: Database Design Spring 2018  
First Midterm Exam**

There are 14 questions and 10 pages in this exam (including this page). To receive credit for a question, answer it according to the instructions given. You can receive partial credit on questions that ask you to write SQL. You have 70 minutes to answer the questions (we will not start until 5 mins into the class period).

Write your name on this cover sheet (see above) AND at the top of each page of this exam. Also, write your UID in the box below. Some questions may be harder than others. Attack them in the order that allows you to make the most progress. If you find a question ambiguous, be sure to write down any assumptions you make. Be neat. If I can't understand your answer, I can't give you credit! Please write all answers within the box provided by the question (if one is provided).

Don't worry if you don't complete all the questions. You receive 25% credit for any question that you leave **COMPLETELY** blank. This credit goes away once you try to start to answer it. In addition, every student will receive a bonus amount of points based on a complicated formula based on the average number of questions (weighted by point value) left empty across the class.

Grades (leave blank --- this will be filled in by graders):

Question 1:	/1
Question 2:	/2
Question 3:	/3
Question 4:	/2
Question 5:	/3
Question 6:	/3
Question 7	/4
Question 8	/4
Question 9	/4.5
Question 10	/4
Question 11	/2
Question 12	/7
Question 13	/4
Question 14	/6.5
Total:	/50

UID: 114462091

Name: Yufan Huang

This entire midterm will use the following schema for a simplified Piazza/Reddit-like discussion board. The database contains the following tables:

1. Users: user\_id, name, email, password
2. Topics: topic\_id, name, description
3. Moderators: user\_id, topic\_id
4. Posts: post\_id, author\_id, topic\_id, parent\_id, title, content

Created (correctly and without errors) according using the following SQL commands:

```
create table users(user_id int primary key, name varchar(10),  
email varchar(20), password varchar(10));  
  
create table topics(topic_id int primary key, name varchar(10),  
description varchar(20));  
  
create table moderators(user_id int references users,  
topic_id int references topics);  
  
create table posts(post_id int primary key,  
author_id int references users,  
topic_id int references topics, parent_id int,  
title varchar(15), content varchar(100));
```

The references keyword creates a foreign-key/primary-key constraint. Thus, for example, user\_id in the moderators table is a foreign key that references the users table.

The users table contains basic information like name, email, etc., about the users of the system.

The topics table contains a list of topics, e.g., Politics, Weather, Science, etc. Each topic has a unique id, name (e.g., "Politics"), and a short description about it.

Every topic has one or more moderators, which we store in the moderators table.

Users of the system can create new notes about any topic and other users can comment on notes. The posts table contains all the notes and the comments associated with it. For example, user "1" may create note "10" on topic "3". Since it is a note, and not a comment on an existing note, we will set the parent\_id to NULL. Thus, the complete row for this note is:

```
(10, 1, 3, NULL, 'This is n10', 'content of note').
```

If another user "2", comments on note "10", this comment gets a new id (say "11"), and we add it to the posts table, with "10" as the parent\_id. That is, the complete row for this comment is:

```
(11, 2, 3, 10, 'First comment', 'content of comment')
```

For simplicity, assume the topic\_id for all comments of a note is the same as the note itself.

Name: yufan Huang

- Question 1 [1 point]: Were the create table commands given above DDL statements or DML statements?

DDL

- Question 2 [2 points]: Write SQL to insert a new user into the users table, with user\_id of 1, whose name is joe, email is joe@umd.edu, and password is 4321.

insert into users values( 1, 'joe', 'joe@umd.edu', '4321');

- Question 3 [3 points]: Assume other users have been inserted into the users table subsequently to question 2, but thus far, only insert statements have been sent to the database. Write SQL to update the value of joe's password (the tuple that you inserted as part of question 1, and not any other tuple in the users table, no matter what values they may have for any of their attributes) to 1234.

update users set password ('1234') where user\_id = 1;

- Question 4 [2 points]:

Write SQL to delete all users that have an email address of joe@umd.edu

delete from users where email = 'joe@umd.edu';

- Question 5 [3 points]:

Assume (for this question only) that only the users table exists in the database (the other tables haven't been created yet) and that users have already been inserted into it. One way to remove the password attribute from the users table is run the following two SQL statements:

```
drop table users;
create table users(user_id int primary key, name varchar(10),
email varchar(20));
```

which drops the users table, and creates a new one without the password attribute.

True or false: (Assume the users table is not empty.) This solution (shown above) for removing the password attribute causes more data loss than just the password values. In particular, a new table will be created, but all the tuples that had been inserted into the old table will gone.

T

True or false: There's no way to remove the password attribute from the users table (with or without data loss) using a single SQL statement.

F

True or false: There's no way to remove the password attribute from the users table using a single SQL statement without losing more data than just the password values.

F

Name: Yufan Huang

Question 6 [3 points]: Write a simple query that returns the user\_id of all users whose email address ends in edu

```
Select user_id from users where  
email = '%edu';
```

Question 7 [4 points]:

We want to compute the set of users who have contributed a lot of posts (not counting comments --- just original notes) --- more than 100 --- and who are not moderators (for any topic). Modify the WHERE clause in the query below to compute the answer.

```
SELECT p.author_id  
FROM Posts p  
WHERE <your-solution-will-be-inserted-here>  
GROUP BY p.author_id  
HAVING COUNT(*) > 100
```

```
where not exist (select * from moderators where  
user_id = p.author_id)
```

Name: *yufan Huang*

Question 8 [4 points]:

Define the **mode user\_id** (there could be more than one) in the posts table as the user(s) with the most notes or comments in the posts table (in other words, the user\_id(s) that appear(s) the most frequently posts table). Circle all of the following queries that correctly return an ordered list of post\_ids of all the posts by the mode user\_id(s) in the posts table. Just circle the first line (Choice X), not the rest of the query.

**Choice A:**

```
with user_posts as
    (select author_id, post_id, count(*) as ct
     from posts group by author_id)
select post_id from user_posts
where ct = (select max(ct) from user_posts) order by post_id;
```

**Choice B:**

```
with user_posts as
    (select author_id, count(*) as ct
     from posts group by author_id)
select post_id from user_posts natural join posts
where ct = (select max(ct) from user_posts) order by post_id;
```

**Choice C:**

```
with user_posts as
    (select author_id, count(*) as ct
     from posts group by author_id),
mode_users as
    (select author_id from user_posts
     where ct = (select max(ct)
                  from user_posts))
select post_id from posts
where author_id in (select * from mode_users) order by post_id;
```

**Choice D:**

```
with user_posts as
    (select author_id, count(*) as ct
     from posts group by author_id),
mode_users as
    (select author_id from user_posts
     where ct = (select max(ct) from user_posts))
select post_id
from posts inner join mode_users
      on (posts.author_id = mode_users.author_id)
order by post_id;
```

Question 9 [4.5 points]: Circle each of the following SQL queries on the next page that will correctly return, for every user in the database, the number of comments (not including original notes) that the user has authored. (Hint: at least one of the queries is correct).

Name: Yufan Huang

Please note that the 9 queries below are in groups of three, where each group of three is the same exact query, except the type of join is changed (left outer join, right outer join, and full outer join). Just circle the first line (Choice X), not the rest of the query.

**Choice A:**

```
SELECT u.user_id, COUNT(*) AS number_of_comments
FROM users AS u LEFT OUTER JOIN posts AS p ON u.user_id = p.author_id
WHERE p.parent_id IS NOT NULL
GROUP BY u.user_id
```

**Choice B:**

```
SELECT u.user_id, COUNT(*) AS number_of_comments
FROM users AS u RIGHT OUTER JOIN posts AS p ON u.user_id = p.author_id
WHERE p.parent_id IS NOT NULL
GROUP BY u.user_id
```

**Choice C:**

```
SELECT u.user_id, COUNT(*) AS number_of_comments
FROM users AS u FULL OUTER JOIN posts AS p ON u.user_id = p.author_id
WHERE p.parent_id IS NOT NULL
GROUP BY u.user_id
```

**Choice D:**

```
SELECT u.user_id, count(p.parent_id)
FROM users as u LEFT OUTER JOIN posts as p on u.user_id = p.author_id
GROUP BY u.user_id;
```

**Choice E:**

```
SELECT u.user_id, count(p.parent_id)
FROM users as u RIGHT OUTER JOIN posts as p on u.user_id = p.author_id
GROUP BY u.user_id;
```

**Choice F:**

```
SELECT u.user_id, count(p.parent_id)
FROM users as u FULL OUTER JOIN posts as p on u.user_id = p.author_id
GROUP BY u.user_id;
```

**Choice G:**

```
SELECT u.user_id, count(distinct p.parent_id)
FROM users as u LEFT OUTER JOIN posts as p on u.user_id = p.author_id
GROUP BY u.user_id;
```

**Choice H:**

```
SELECT u.user_id, count(distinct p.parent_id)
FROM users as u RIGHT OUTER JOIN posts as p on u.user_id = p.author_id
GROUP BY u.user_id;
```

**Choice I:**

```
SELECT u.user_id, count(distinct p.parent_id)
FROM users as u FULL OUTER JOIN posts as p on u.user_id = p.author_id
GROUP BY u.user_id;
```

Name: YM fan Huang

Question 10 [4 points]: It is possible that some topics have exactly one moderator (every topic will have at least one moderator). It is also possible that a user has authored notes and/or comments on the same topic for which that user is the only moderator. We want to identify such notes and comments, so that we can add more moderators to that topic. Specifically, we wish to find the set of all notes/comments where the author of the note/comment is the sole moderator of the post's topic. Complete the WHERE clause in the query below. (Note: the use of the sum aggregator below is not a typo.)

```
WITH topics_with_one_mod AS (
    SELECT topic_id, sum(user_id) as uid
    FROM moderators
    GROUP BY topic_id
    HAVING COUNT(*) = 1)
SELECT p.post_id
FROM posts AS p
WHERE EXISTS <write solution here>;
```

```
Select * from topics_with_one_mod p1 where
p1.topic_id = p.topic_id
```

Question 11 [2 points]: True or false: No matter what the contents are of the moderators and users tables, it is **impossible** for the following query to return more than 0 rows:

```
SELECT * FROM users NATURAL FULL OUTER JOIN moderators
EXCEPT
SELECT * FROM users NATURAL LEFT OUTER JOIN moderators
```

T

Name:

yufan Huang

Question 12: After loading data into the database, the user ran the following query:  
select \* from moderators;

and got the response:

user_id	topic_id
1	1
2	2
2	1
2	3

1 - 1  
2 < 2  
\\ 1  
3

How many rows are returned by the following queries? (If any of them will return an error, write "error".)

[1 point]

SELECT max(user\_id) FROM moderators;

error

[1 point]

SELECT DISTINCT user\_id FROM moderators;

2

[1 point]

SELECT user\_id, topic\_id, count(\*)  
FROM moderators  
GROUP BY user\_id, topic\_id;

2

[2 points]

SELECT \* FROM moderators as m1 WHERE m1.user\_id =  
(SELECT m2.user\_id  
FROM moderators as m2  
WHERE m2.user\_id > ALL  
(SELECT m3.user\_id FROM moderators as m3));

0

[1 point]

$\sigma_{(topic\_id = user\_id)}$ (moderators)

2

1 > 1  
1 > 2  
2  
2

[1 point]

$\Pi_{(topic\_id)}$ (moderators)

2

Name:

yutan huang

Question 13 [4 points]: We decide to have a rule that all users' names must begin with the letter 'j'. Otherwise they are not allowed to exist in our database. To enforce this rule, we decide to check every new tuple that is inserted to make sure it meets our standards (that the name must begin with 'j'), and also check every tuple that is updated to ensure that the name wasn't changed to a new value that doesn't meet our standards. I wrote the below trigger in Postgres called `check_name()` that checks the new value of a record (that was either inserted or modified) to ensure that the name attribute starts with 'j'. If it does, then the new value is returned and the insert or modification can proceed. If it does not start with 'j', then null is returned, which causes the insert or modification to fail. You can assume that the code in the function works correctly.

```
CREATE OR REPLACE FUNCTION check_name() RETURNS trigger AS $$  
BEGIN  
    IF NEW.name LIKE 'j%' THEN --if name starts with j  
        RETURN NEW; --return the new tuple to be inserted  
    ELSE  
        RETURN NULL; --reject new tuple because didn't start with j  
    END IF;  
END;  
$$ LANGUAGE plpgsql;
```

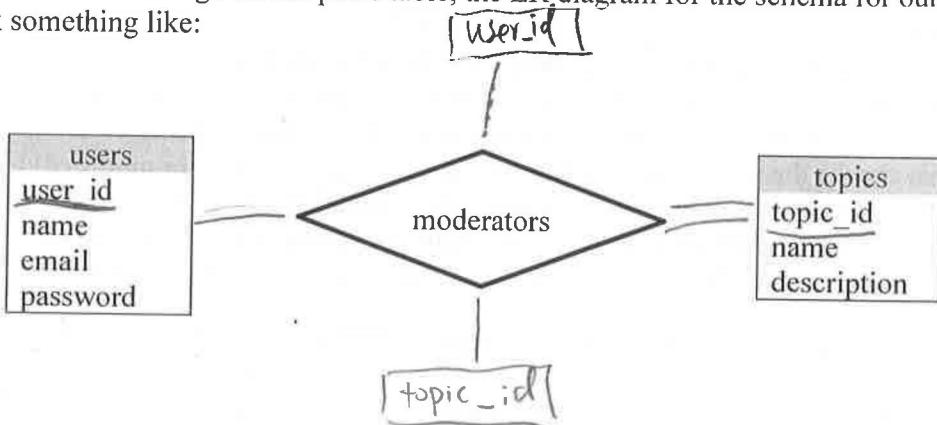
I then ran a `CREATE TRIGGER` command in Postgres to ensure that trigger is fired on the `users` table. However, I removed the part of the statement that says when the trigger should be fired. Please complete the missing part of the `CREATE TRIGGER` statement to tell Postgres when the trigger should be fired.

```
CREATE TRIGGER check_name <your-answer-will-go-here> ON users  
    FOR EACH ROW EXECUTE PROCEDURE check_name();
```

before update or insert

Name: Wfan Huang

Question 14: If we ignore the posts table, the ER diagram for the schema for our database will look something like:



Part 1 [4 points]: Assume the only constraints that exist in the application are those mentioned in the description page on page 2 of this midterm (including the create table statements which define the foreign key constraints). Remember we said that that every topic has at least one moderator. Draw the lines that connect the users and topics entity sets to the moderators relationship set in a way that shows mapping cardinality constraints using the standard notation (from the textbook and lecture notes) of using directed lines that point to the “one” side of a many-to-one, one-to-many, or a one-to-one relationship, and undirected lines for many-to-many relationships. Distinguish between partial and total participation constraints using the standard notation (from the textbook and lecture notes) of single vs. double lines.

Part 2 [1.5 point]: Add any attributes to the moderators relationship set to the figure above (if appropriate).

Part 3 [1 point]: Underline all attributes (either the ones that existed already or the ones you added in Part 2) in the ER diagram above that compose part of a primary key.

