

Reading Homework 11 (Due Wednesday 04/04, 3:20PM)

Due Apr 4, 2018 at 3:25pm**Points** 10**Questions** 8**Available** until May 17, 2018 at 11:59am**Time Limit** None

Instructions

The reading assignment here is: Chapters 12.1, 12.2, 12.3.1.

This quiz was locked May 17, 2018 at 11:59am.

Attempt History

	Attempt	Time	Score
LATEST	Attempt 1	less than 1 minute	0 out of 10

Score for this quiz: **0** out of 10

Submitted Apr 4, 2018 at 3:29pm

This attempt took less than 1 minute.

Question 1

1 / 1 pts

There are usually many different ways to execute the same query.

Correct!☒ True☐ False

Question 2**1 / 1 pts**

Optimizers generally try to minimize the response time for a query.

☐ True☒ False**Correct!****Question 3****1 / 1 pts**

The query optimizer is typically responsible for confirming that the query syntax is valid.

☐ True☒ False**Correct!****Question 4****1 / 1 pts**

Given two plans: (1) a plan that makes 1 random seek and reads 1000 blocks in total,
vs (2) a plan that makes 10 random seeks and reads 10 blocks in total, plan
(1) is better
to use. Use the values from the book: $t_S = 4\text{ms}$ and $t_T = 0.1\text{ms}$.

☐ True☒ False**Correct!**

Question 5**1 / 1 pts**

Given two plans: (1) a plan that makes 100 random seeks and reads 100000 blocks in total, vs (2) a plan that makes 1000 random seeks and reads 1000 blocks in total, plan (1) is better to use. Use the values: $t_S = 5\text{ms}$ and $t_T = 0.01\text{ms}$.

Correct!☒ True☐ False**Question 6****1 / 1 pts**

What would be value of t_T if you assume 8KB blocks and a 400MB/s disk transfer rate?

Correct!☐ 0.01ms☒ 0.02ms☐ 0.05ms☐ 0.2ms**Question 7****2 / 2 pts**

Say we have a primary B+-tree of height 4 on attribute "zipcode" in an "person" table (so there will be many records with the same zipcode, but the relation is sorted by zipcode). Consider a query to find all people in a specific zipcode, and let's say there are a 1000 records with that zipcode. Further, let's say a single relation block

can hold

10 records. Estimate the cost of executing this query (using Figure 12.3).

Assume $t_S = 4\text{ms}$, and $t_T = 0.1\text{ms}$.

Correct!

☐ Approximately 3ms

☒ Approximately 30ms

☐ Approximately 60ms

☐ Approximately 300ms

Question 8

2 / 2 pts

Do the same but with the assumption that the index is a secondary index (i.e., the relation data is not sorted by zipcode).

Correct!

☐ approximately 200ms

☒ approximately 2s

☐ approximately 20s

☐ approximately 200s

Quiz Score: **0** out of 10

This quiz score has been manually adjusted by -10.0 points.