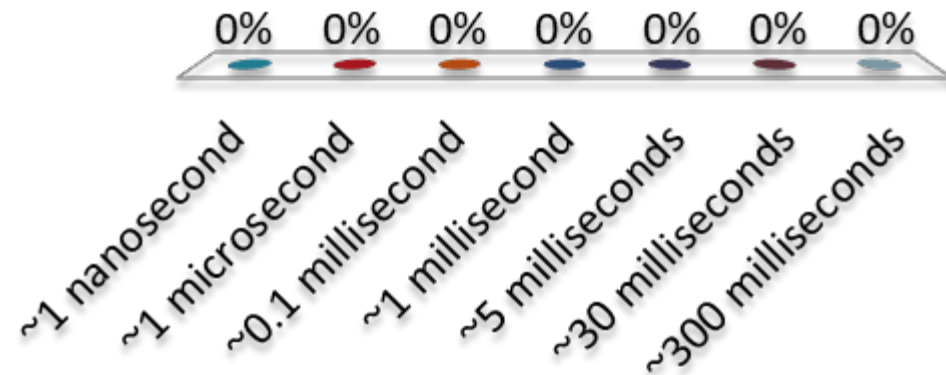


# Sequential File Organization

- ▶ Keep sorted by some search key
- ▶ Insertion
  - Find the block in which the tuple should be
  - If there is free space, insert it
  - Otherwise, must create *overflow pages*
- ▶ Deletions
  - Delete and keep the free space
  - Databases tend to be insert heavy, so free space gets used fast
- ▶ Can become *fragmented*
  - Must reorganize once in a while

Find a record with ID = x in sequential file with 1,000,000,000 records, each record is 100 bytes?

- A. ~1 nanosecond
- B. ~1 microsecond
- C. ~0.1 millisecond
- D. ~1 millisecond
- E. ~5 milliseconds
- F. ~30 milliseconds
- G. ~300 milliseconds



# Sequential File Organization

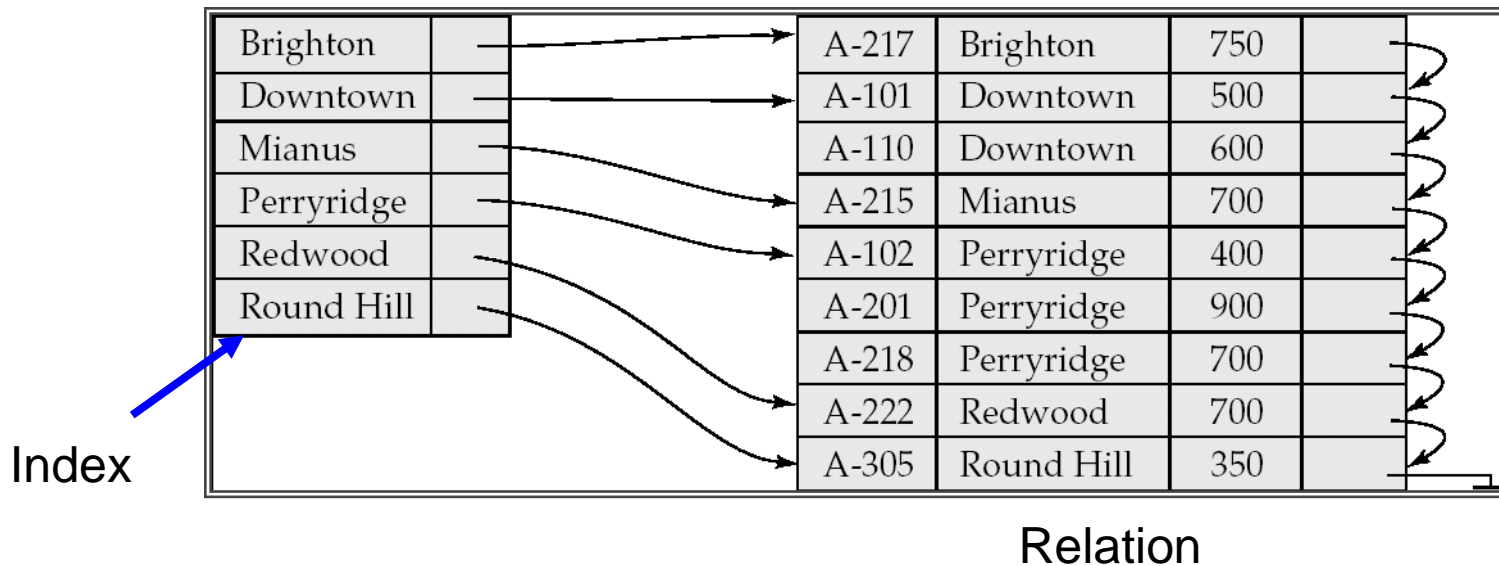
- ▶ What if I want to find a particular record by value ?
  - *Account info for ID = 123*
- ▶ Binary search
  - Takes  $\log(n)$  number of disk accesses
    - Random accesses
  - Too much
    - $n = 1,000,000,000$  --  $\log(n) = 30$
    - Recall each random access approx 10 ms
    - 300 ms to find just one account information
    - $< 4$  requests satisfied per second

# Index

- ▶ A data structure for efficient search through large databaess
- ▶ Two key ideas:
  - The records are mapped to the disk blocks in specific ways
    - Sorted, or hash-based
  - Auxiliary data structures are maintained that allow quick search
- ▶ Think library index/catalogue
- ▶ Search key:
  - Attribute or set of attributes used to look up records
  - E.g. ID for a person table
- ▶ Two types of indexes
  - *Ordered* indexes
  - *Hash-based* indexes

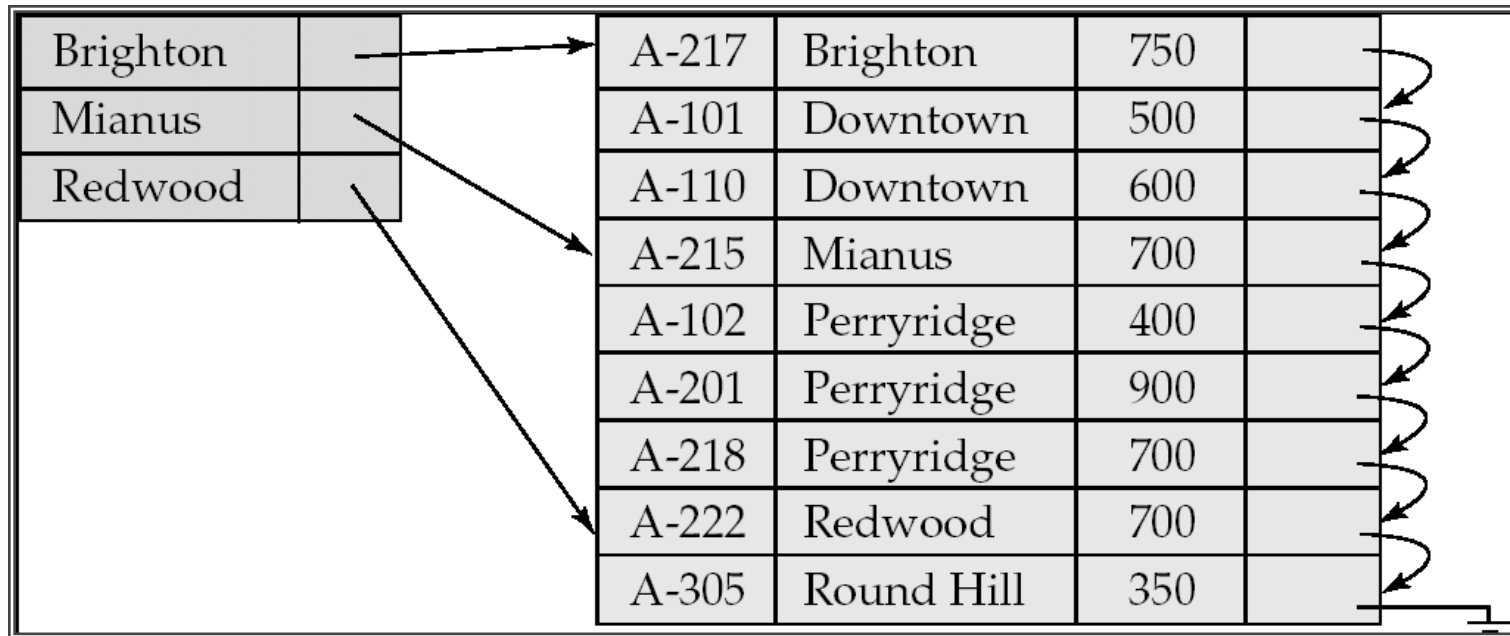
# Ordered Indexes

- ▶ Primary index
  - The relation is sorted on the search key of the index
- ▶ Secondary index
  - It is not
- ▶ Can have only one primary index on a relation




# Primary Sparse Index

- ▶ Every key doesn't have to appear in the index
- ▶ Allows for very small indexes
  - Tradeoff: Must access the relation file even if the record is not present
  - Tradeoff: Must scan through multiple records to find desired record
    - But at least this scan is sequential




How many index keys / records will be accessed to find X which happens to be the 100<sup>th</sup> record of a table, when a sparse index contains an entry for every 40 records (i.e. it has pointers to the 1<sup>st</sup>, 41<sup>st</sup>, 81<sup>st</sup>, 121<sup>st</sup>, etc. records)

- A. 1
- B. 20
- C. 24
- D. 40
- E. 41
- F. 100



How many index keys / records will be accessed to find X which happens to be the 100<sup>th</sup> record of a table, when a dense index is used instead of a sparse index?


- A. 1
- B. 20
- C. 24
- D. 40
- E. 41
- F. 100





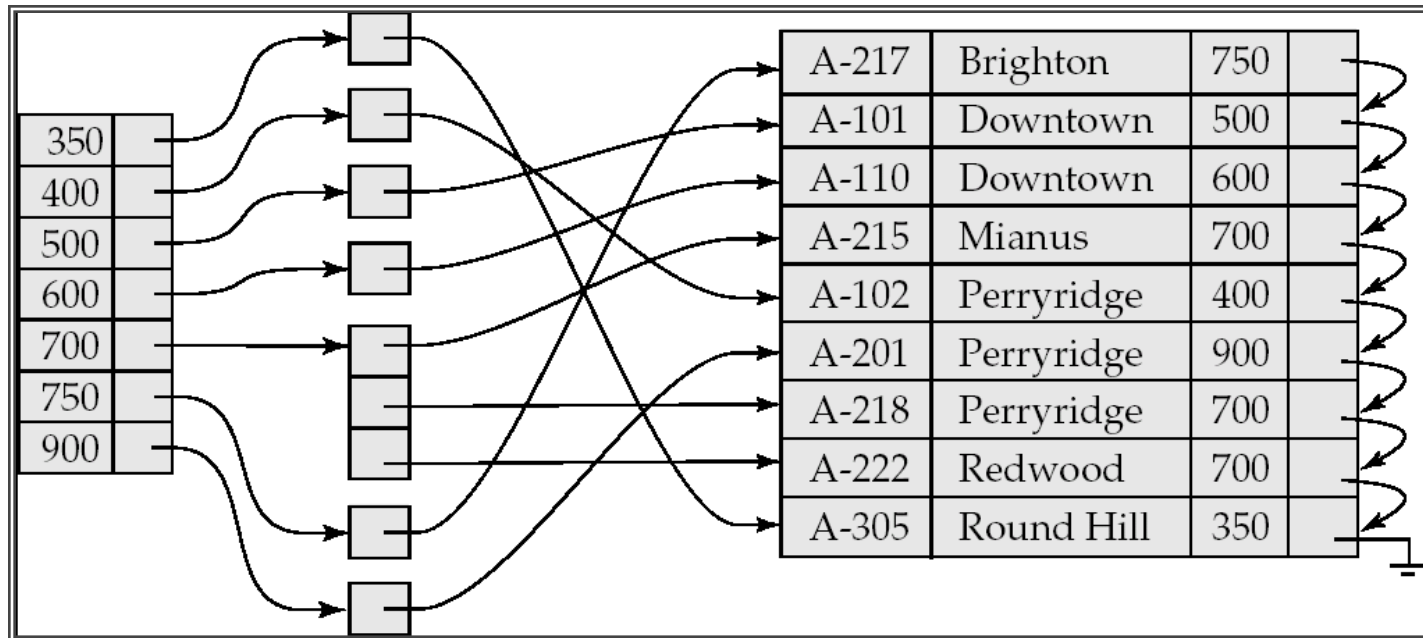
# When should a dense index be used instead of a sparse index (assume data is sorted by indexed attribute)

- A. Never
- B. When we want to find one of the first few tuples in a table
- C. When we want to find one of the last few tuples in a table
- D. When the entire dense index fits in memory
- E. When random IO is just as slow as sequential IO



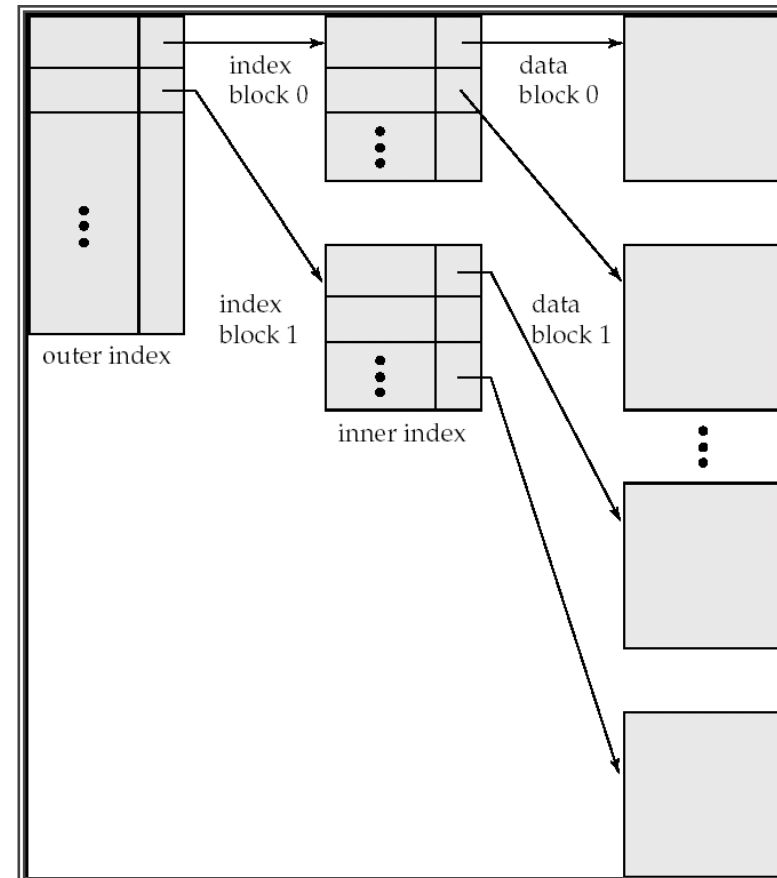
# Secondary Index

- ▶ Relation sorted on *branch*
- ▶ But we want an index on *balance*
- ▶ Must be dense
  - Every search key must appear in the index




# Multi-level Indexes

- ▶ What if the index itself is too big for memory ?
- ▶ Relation size =  $n = 1,000,000,000$
- ▶ Block size = 100 tuples per block
- ▶ So, number of pages = 10,000,000
- ▶ Solution
  - Build a **sparse outer** index on the index itself



# What do you think is the biggest downside of multi-level indexes from the previous slide?

- A. They take too much space!
- B. They only work for primary indexes (when data is sorted by the indexed attribute)!
- C. There is one random IO per level in the hierarchy!
- D. Their performance will be bad in the face of many inserts!



# Multi-level Indexes

- ▶ What about keeping the index up-to-date ?
  - Tuple insertions and deletions
    - This is a static structure
    - Need overflow pages to deal with insertions
  - Works well if no inserts/deletes
  - Not so good when inserts and deletes are common
- ▶ This is the main motivation for B-trees