

CMSC 424: Database Design

Spring 2018

Introductory Material



Database Systems

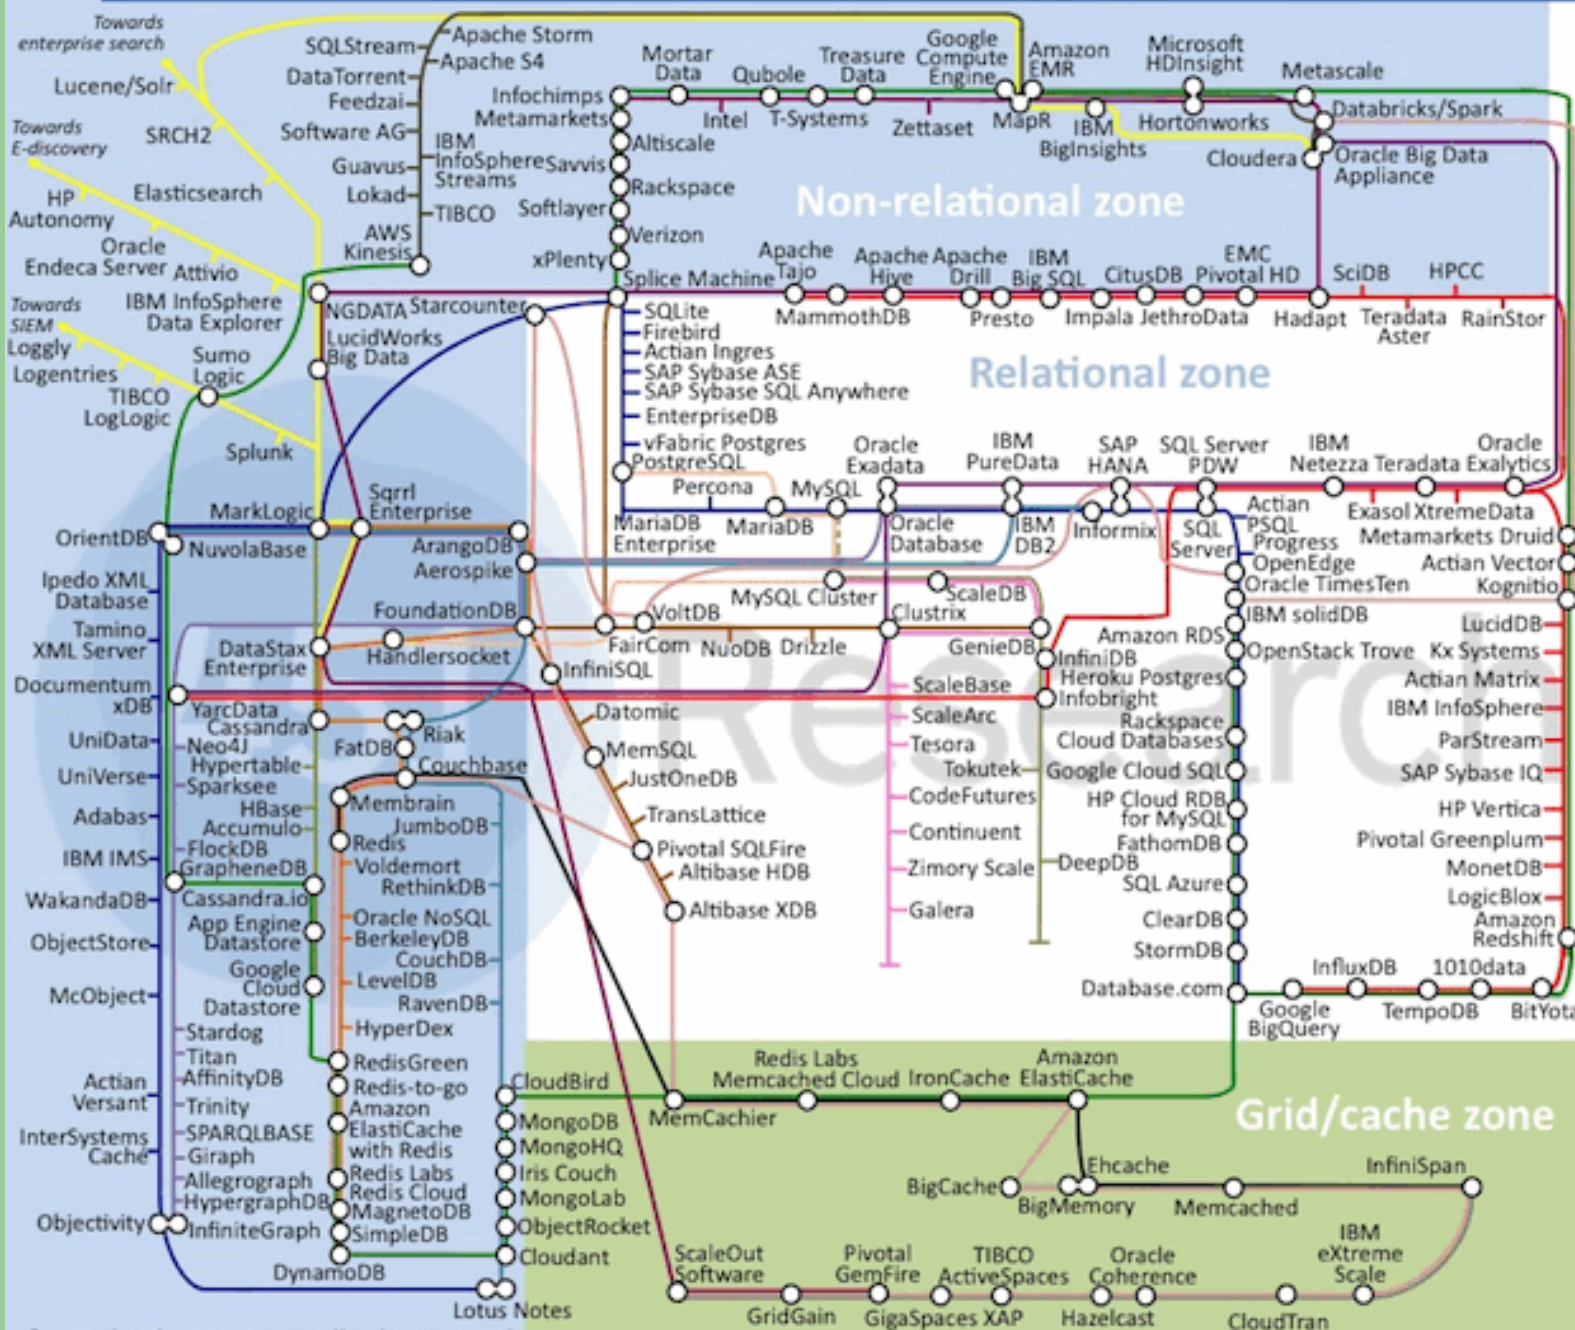
- Database: a structured, queryable collection of records
- Database management system (DBMS): software to manage and query the database
 - 20 billion dollar industry
 - Deployed everywhere
 - At the low end, you can get a DBMS for free; at the high end, they can go for 7 figures

Data Platforms Landscape Map – February 2014

451 Research

Key:

- General purpose
- Specialist analytic
- -as-a-Service
- - - NoSQL extension
- BigTables
- Graph
- Document
- Key value stores
- Key value direct access
- Hadoop
- - - NewSQL extension
- MySQL storage engine
- Advanced clustering/sharding
- New SQL databases
- Data caching
- Data grid
- Search
- Appliances
- Off-heap memory
- In-memory
- Stream processing



Overview

- This course covers
 - How to use a DBMS
 - If you build an app, you will likely use a DB
 - How to build a DBMS
 - Database systems are highly concurrent, complex systems
 - Learning how they work:
 - Will make you a better programmer
 - Will open up more career opportunities

Prerequisites

- Minimum grade of C- in CMSC351 and CMSC330

Readings

- Textbook: “Database System Concepts”
 - Silberschatz, Korth & Sudarshan; 6th Edition, McGraw Hill.
- Readings listed on syllabus refer to 6th edition.
 - Earlier edition OK, but you have to figure out the mapping.
- Some other material linked to from the syllabus

Exams

- 2 Midterms
 - March 5 and April 9 in class
- Final exam on May 15 from 1:30-3:30PM

Projects / Programming Assignments

- Done individually
- You may ask, and answer, **general** questions on Piazza
- Assignments will use a mixture of languages, with Python and Java being the most common.

Clickers / In-class participation

- We will be using clickers this semester
- Your answers will count toward final grade
- Any question that majority of class gets wrong will be deemed unfair and thrown out
- You can avoid answering questions for any three classes of your choosing
 - As soon as you answer the first question from a class, you can no longer count it as a missed class

Reading homeworks

- Online quizzes before class
- No late submissions accepted
 - But you can skip 2 for any reason

Course grade

- Projects/programming assignments: 30%
- Reading homeworks: 5%
- Mid-term exams: 32%
- Final Exam: 25%
- Class participation (clicker answers): 8%

Teaching Fellows

- Amit Chavan - Wed 9AM-noon
- Hui Miao - Thu 10AM-1PM
- Souvik Bhattacharjee - Tue 9AM-noon
- Virinchi Srinivas - Mon 9AM-noon
- Adam Ackerman - Mon noon - 1PM; Friday 11:30AM-1:30PM
- Ayokiitan Akalaa - Mon/Wed 2-3:30PM
- Kevin Wittmer - Mon/Wed 5:30-7:00PM
- Brian Mitchell - Thu 10:30AM - noon, Fri 10AM-11:30AM
- Anthony Pingelli - Mon 1PM-2PM, Wed noon-2PM

- Bottom line:
 - M/W Office hours 9AM-class time, 5:30-7PM
 - Tu/Th Office hours morning/early afternoon
 - Fri office hours 10AM-1:30PM

Assignment 0

- Assignment 0 (setting up your programming environment) will be released shortly
 - Vagrant+VirtualBox, PostgreSQL, IPython, etc.
- TA office hours start tomorrow in case you need help

True story

- Needed to create an online RSVP application
- Specs:
 - For a series of meetings
 - Meeting would only be held if $\geq X$ people committed to attending
 - Before each meeting, send out an e-mail, asking people to RSVP if they can come
 - Count responses, then decide if meeting is on

First Cut

- Have everyone register an account with some basic information
 - Form-based Web-page writing to a file (accounts.txt) in the backend (Perl)
- E-mail points to a different Web-form for people to sign-in and say 'yes' or 'no'
 - Again, results written to file (meeting01-12.txt)
- Cron-script that periodically counts how many responses and either
 - Sends out another reminder e-mail OR
 - Makes final decision
 - Status file (decided.txt) edited to stop mind-changing

End of story?

- People are cranky – hate e-mails
 - “I never can make Monday meetings”
- Not everyone should be counted
- “yes”, “no” too restrictive
- Problem: implementation nightmare
 - Everything requires custom code
 - Bugs crept up through lack of data independence

It gets worse

- Concurrency became a problem
 - Need to lock entire file to edit just one line
 - In some cases, concurrency resulted in miscounts

And worse ...

- Performance became an issue with scale
 - find the e-mails of all people coming (from ID)
 - view status of all upcoming meetings

Lessons Learned

- Never tell anyone you know how to create a Web application (or mobile or whatever)
- There's no such thing as a simple app
- Should have used a DBMS

Should have used a DBMS

- Generic query interface
- Handles data independence
- Handles concurrency
- Handles recovery from crash
- Designed for performance (indexes, address offsetting, result caching)

Relational Data Model

- Data stored in tables, tables have attributes

Relational Model

People

| ID | Name | E-mail | Phone | Gender |
|----|---------|--|--------------|--------|
| 1 | Joe | joe@yale.edu | 555-555-5555 | M |
| 2 | Jill | jill@yale.edu | 555-555-5555 | F |
| 3 | Jasmine | jasmine@yale.edu | 555-555-5555 | F |
| 4 | Jeff | jeff@yale.edu | 555-555-5555 | M |

Meetings

| ID | Date | Time | Location |
|----|--------|--------|----------|
| 1 | 15-Jan | 7:30AM | AKW 500 |
| 2 | 15-Jan | 4:30PM | AKW 500 |
| 3 | 17-Jan | 7:30AM | AKW 500 |

People-Meetings

| Person ID | Meeting ID | Status |
|-----------|------------|--------|
| 1 | 1 | yes |
| 2 | 1 | yes |
| 3 | 1 | no |
| 4 | 1 | yes |
| 1 | 2 | yes |
| 2 | 2 | no |
| 3 | 2 | no |
| 4 | 2 | no |

Other Models Possible

- (Joe, joe@yale.edu, 555-555-5555, M, (((meeting 1, Jan-15, 7:30AM, AKW 500), yes), ((meeting 2, Jan-15, 4:30PM, AKW 500), no))), (Jill, ...)

Database queries

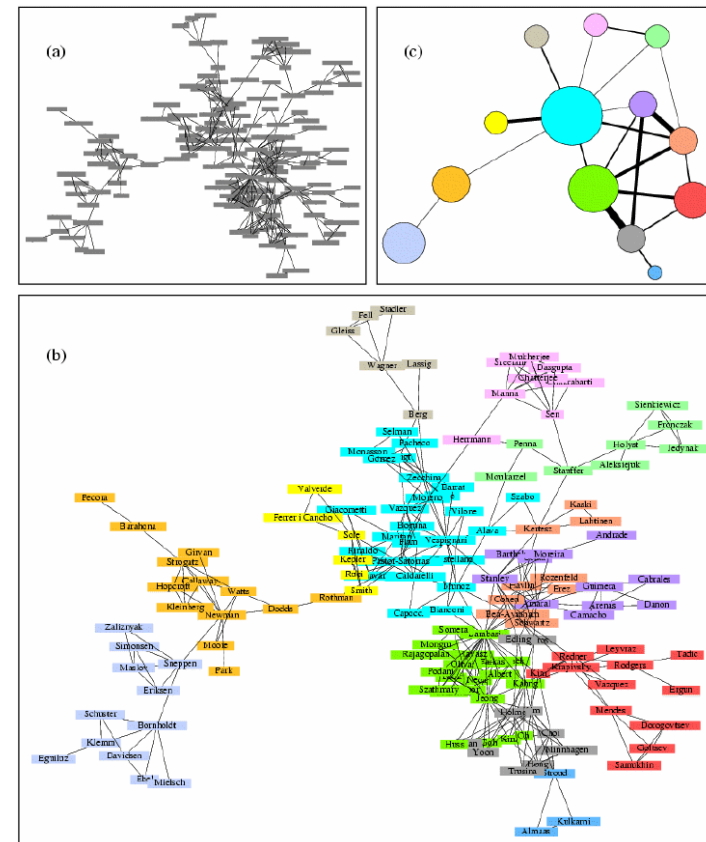
- SELECT e-mail
FROM People
WHERE Gender='F'
- SELECT People.name
FROM People, People-Meetings
WHERE People.ID=People-Meetings.PersonID
AND People-Meetings.Status = 'yes'
AND People-Meetings.MeetingID = 1

Other Constructs

- INSERT, DELETE, UPDATE

Structured vs Unstructured Data


- Some data has a little **more complicated structure**
 - E.g graph structures
 - Map data, social networks data, the web link structure etc
 - Can convert to tabular forms for storage, but may not be optimal
 - Queries often reason about graph structure
 - *Find my “Erdos number”*
 - *Suggest friends based on current friends*
 - Growing importance in recent years in a variety of domains: Biological, social networks, web...



Structured vs Unstructured

- Increasing amount of data in a **semi-structured format**
 - XML/JSON – Self-describing tags (HTML ?)
 - Complicates a lot of things
- A huge amount of data is **unstructured**
 - Books, WWW
 - Amenable to pretty much only text search... so far
 - Information Retrieval research deals with this topic
 - What about Google search ?
 - Google search is mainly successful because it uses the structure (in its original incarnation)
- Video ? Music ?
 - Can represent in DBMS's, but can't really operate on them

```
<Symbol>List</Symbol>
<Function>
  <Symbol>List</Symbol>
  <Symbol>Automatic</Symbol>
  <Number>4.</Number>
</Function>
<Function>
  <Symbol>List</Symbol>
  <Symbol>Automatic</Symbol>
  <Number>6.</Number>
</Function>
</Function>
</Option>
</Options>
</Notebook>
```



circle size == page importance == **pagerank**
more incoming links → higher pagerank
incoming links from important pages → higher pagerank

DBMSs to the Rescue

- Massively successful for *highly structured data*
 - Why ? Structure in the data (if any) can be exploited for ease of use and efficiency
 - How ?
 - Two Key Concepts:
 - Data Modeling: Allows reasoning about the data at a high level
 - e.g. “emails” have “sender”, “receiver”, “...”
 - Once we can describe the data, we can start “querying” it
 - Data Abstraction/Independence:
 - Layer the system so that the users/applications are insulated from the low-level details

What we will cover...

- Representing information
 - data modeling
 - semantic constraints
- Languages and systems for querying data
 - complex queries & query semantics
 - over massive data sets
- Concurrency control for data manipulation
 - ensuring transactional semantics
- Reliable data storage
 - maintain data semantics even if you pull the plug
 - fault tolerance

What we will cover...

- Representing information
 - data modeling: *relational models, E/R models, XML/JSON*
 - semantic constraints: *integrity constraints, triggers*
- Languages and systems for querying data
 - complex queries & query semantics: *SQL*
 - over massive data sets: *indexes, query processing, optimization, parallelization/cluster processing, streaming*
- Concurrency control for data manipulation
 - ensuring transactional semantics: *ACID properties, distributed consistency*
- Reliable data storage
 - maintain data semantics even if you pull the plug: *durability*
 - fault tolerance: *RAID*

For next class

- Read Chapter 2 of textbook
 - (Also a good idea to skim chapter 1)