

Query Optimization

Query Optimization

- Why ?

- Many different ways of executing a given query
- Huge differences in cost

- Example:

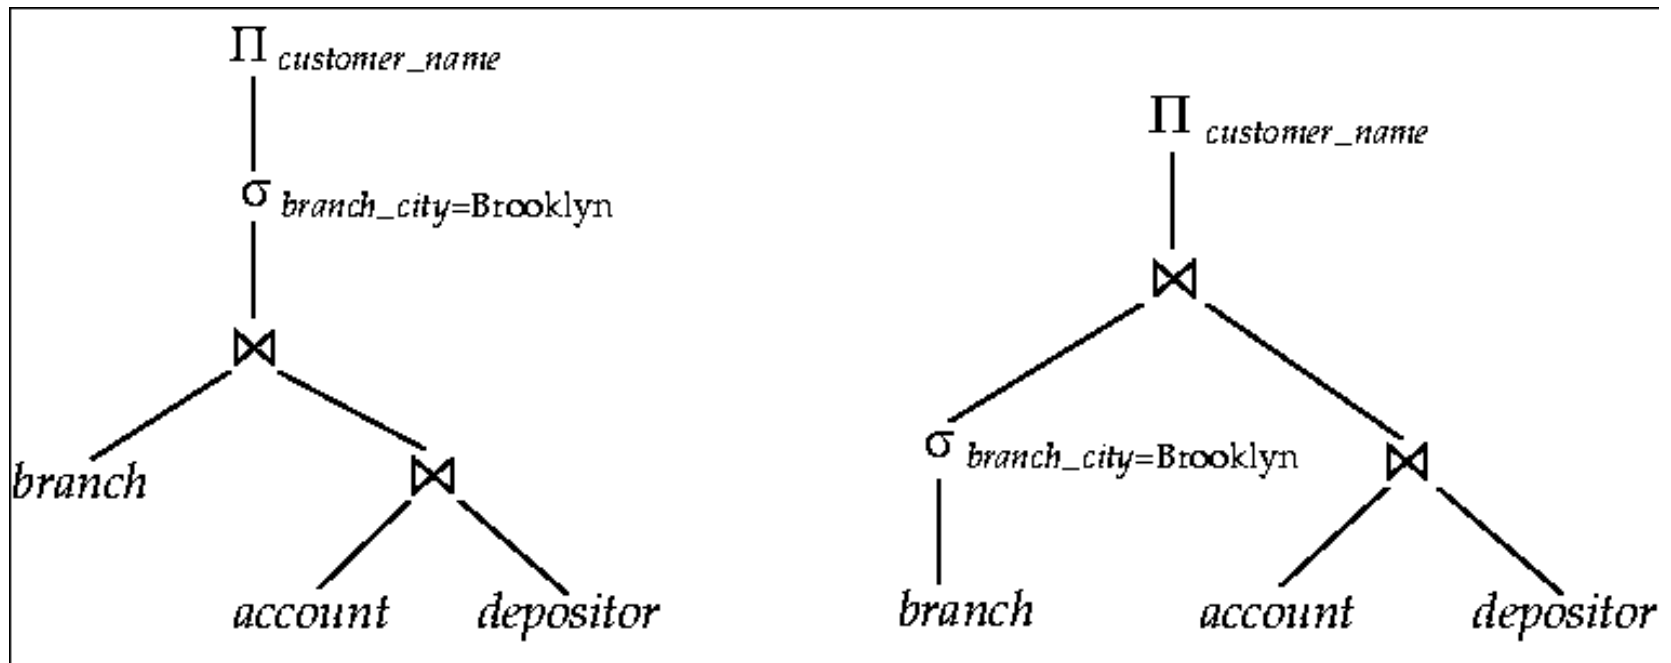
- `select * from person where ssn = "123"`
- Size of person = 1GB
- Sequential Scan:
 - Takes $1\text{GB} / (100\text{MB/s}) = 10\text{s}$
- Use an index on SSN (assuming one exists):
 - Approx 4 Random I/Os = 20ms

Query Optimization

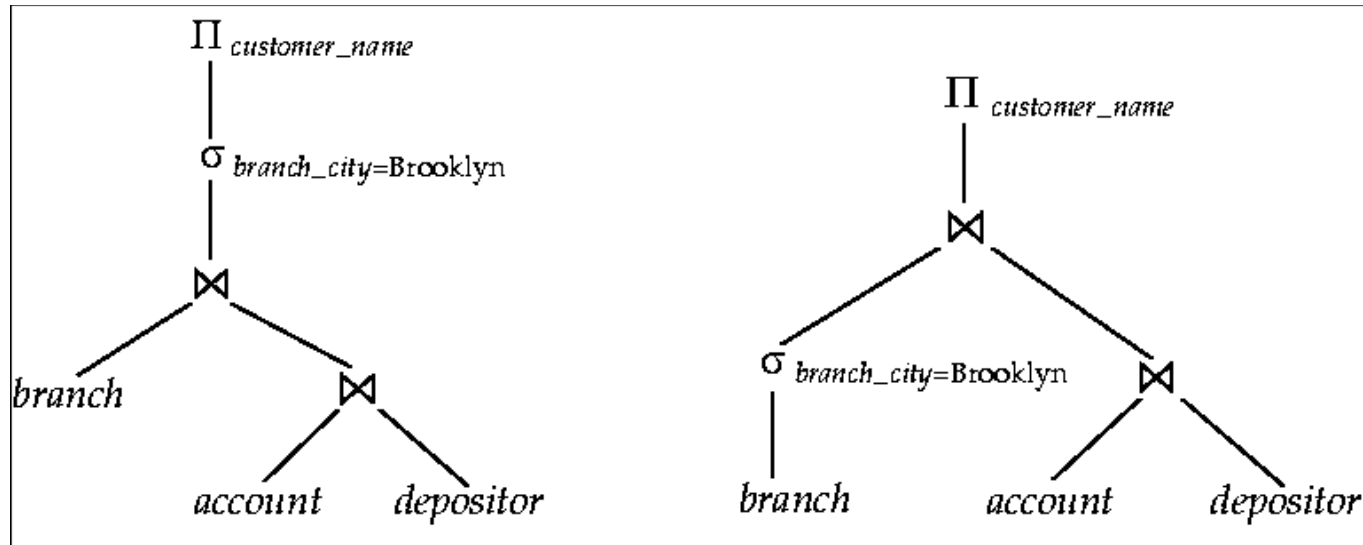
- Many choices
 - Using indexes or not, which join method (hash, vs merge, vs NL)
 - What join order ?
 - Given a join query on R, S, T, should I join R with S first, or S with T first ?
- This is an optimization problem
 - Number of different choices is very very large
 - Step 1: Figuring out the solution space
 - Step 2: Finding algorithms/heuristics to search through the solution space

Query Optimization

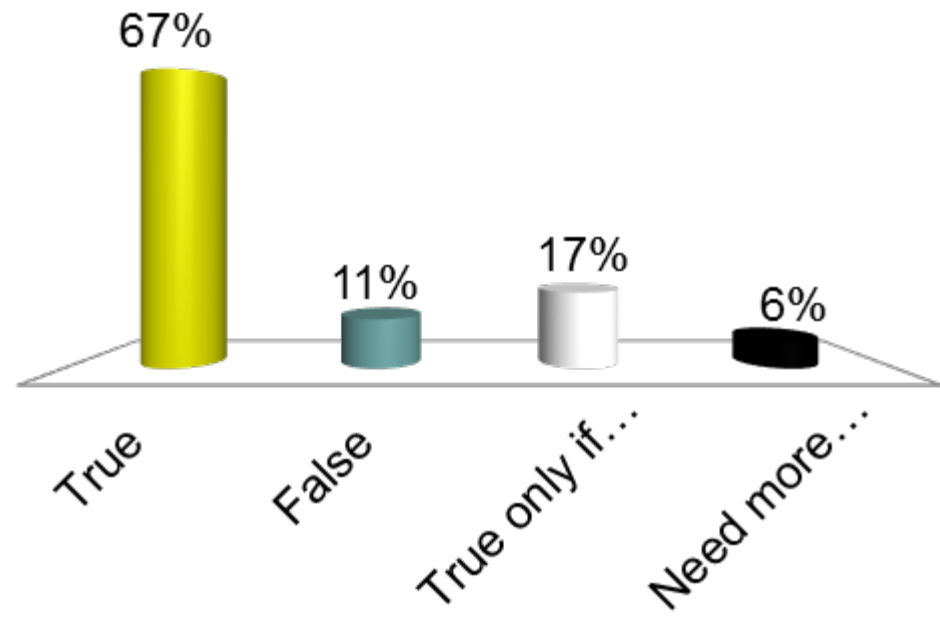
- Relational expressions
 - Drawn as a tree
 - List the operations and the order



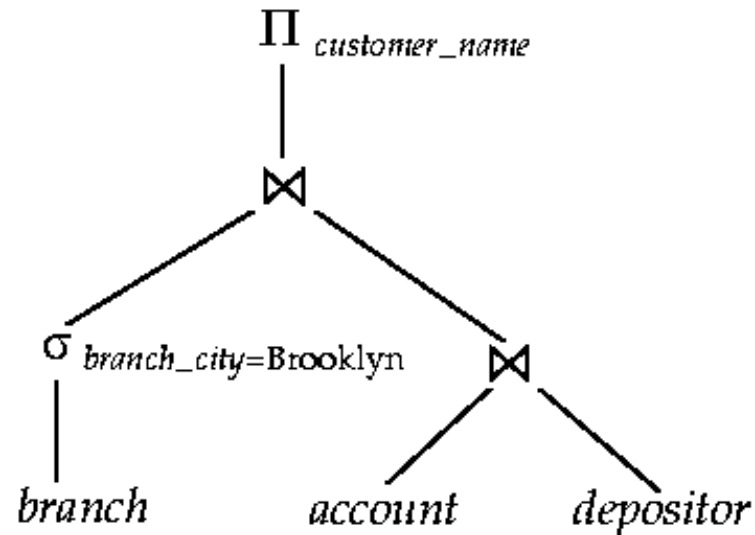

These query plans produce the same final result



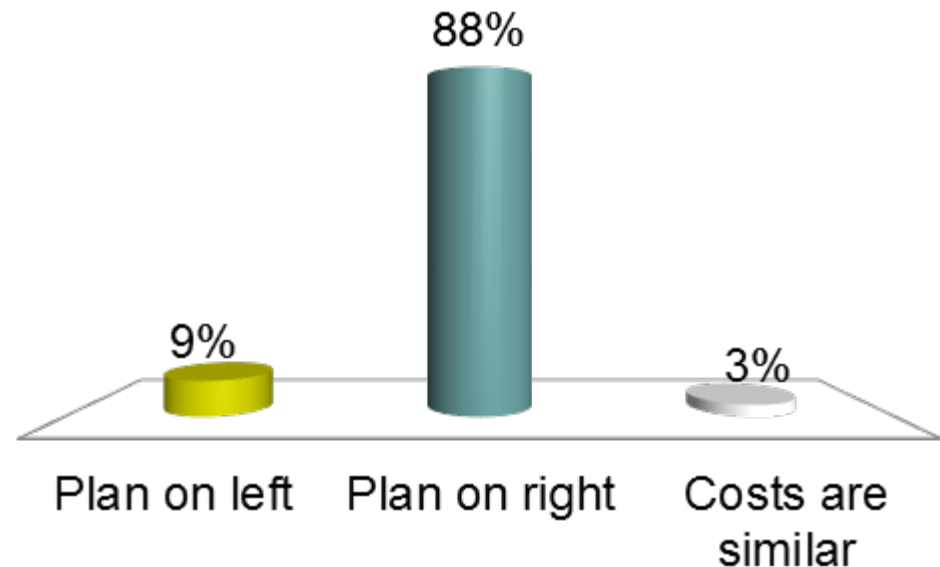
- A. True
- B. False
- C. True only if branch_city cannot be NULL
- D. Need more information



Which query plan is probably faster?

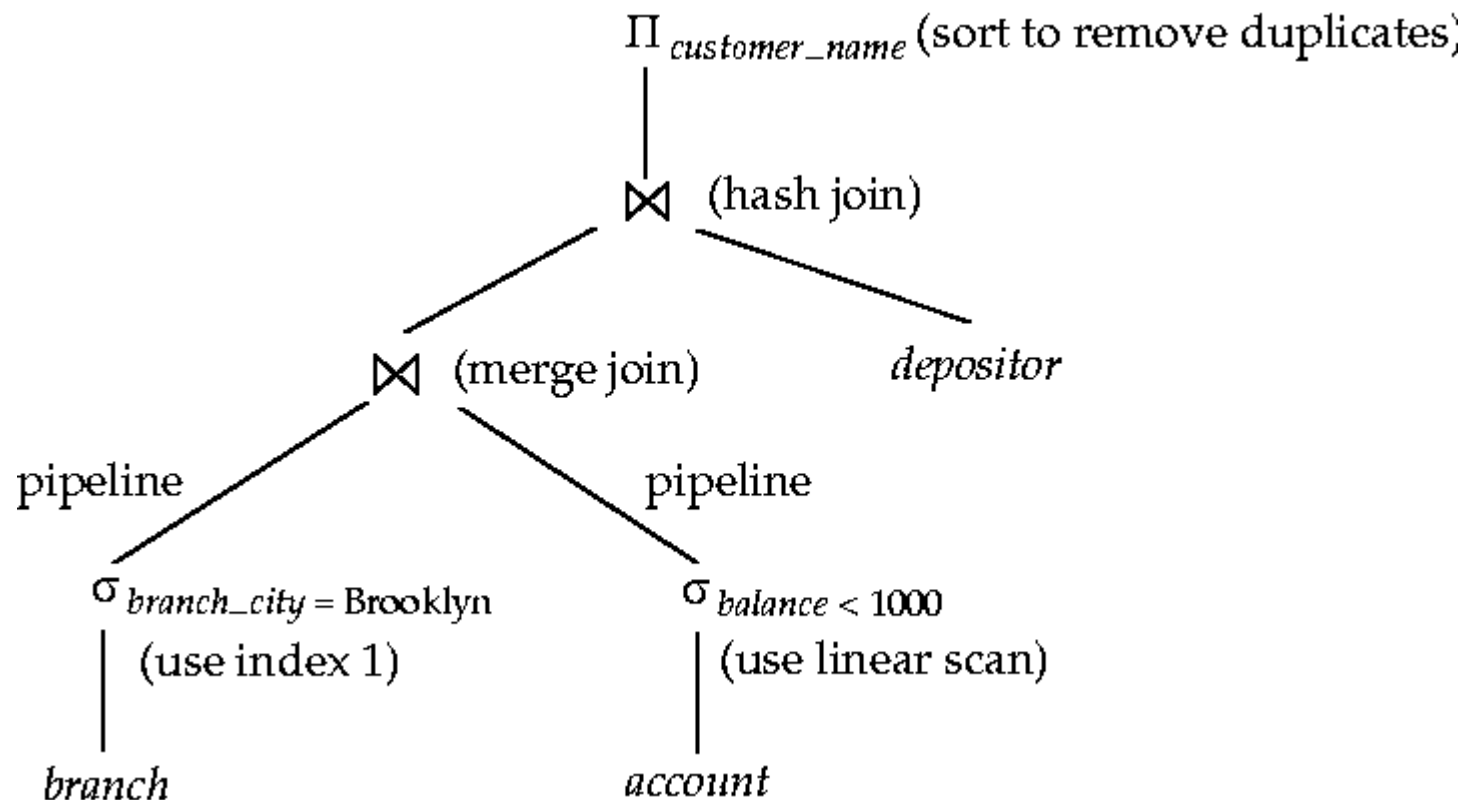


- A. Plan on left
- B. Plan on right**
- c. Costs are similar



Query Optimization

- Execution plans
 - Evaluation expressions annotated with the methods used



Query Optimization

- Steps:
 - Generate all possible execution plans for the query
 - Figure out the cost for each of them
 - Choose the best
- Not done exactly as listed above
 - Too many different execution plans for that
 - Typically interleave all of these into a single efficient search algorithm

Query Optimization

- Steps (detail):

- Generate all possible execution plans for the query

- First generate all equivalent expressions
- Then consider all annotations for the operations

- Figure out the cost for each of them

- Compute cost for each operation

- Using the formulas discussed in prior weeks
- One problem: How do we know the number of result tuples for, say,
 ▪ Count them or estimate...

$\sigma_{balance < 2500}(account)$

- Choose the best

Cost estimation

- Computing operator costs requires information like:
 - Primary key ?
 - Sorted or not, which attribute
 - So we can decide whether need to sort again
 - How many tuples in the relation, how many blocks ?
 - How many tuples match a predicate like “age > 40” ?
 - E.g. Need to know how many index pages need to be read
 - Intermediate result sizes
 - E.g. (R JOIN S) is input to another join operation – need to know if it fits in memory
 - And so on...

Cost estimation

- Some info is static and maintained in the metadata
 - Primary key ?
 - Sorted or not, which attribute
 - So we can decide whether need to sort again
 - How many tuples in the relation, how many blocks ?
- Typically kept in special “catalog” table in the DB
 - E.g., “all_tab_columns” in Oracle
 - E.g., “pg_database” in Postgres

Cost estimation

- Others need to be estimated:
 - How many tuples match a predicate like “age > 40” ?
 - Intermediate result sizes
- The problem variously called:
 - “intermediate result size estimation”
 - “selectivity estimation”
- Very important to estimate reasonably well
 - e.g. consider “SELECT * FROM R WHERE zipcode = 20742”
 - We estimate that there are 10 matches, and choose to use a secondary index (remember: random I/Os)
 - Turns out there are 10000 matches
 - Using a secondary index very bad idea
 - Optimizer often chooses nested-loop joins if one relation very small... underestimation can be very bad

Selectivity Estimation

- Basic idea:


- Maintain some information about the tables
 - More information → more accurate estimation
 - More information → higher storage cost, higher update cost
- Make uniformity and randomness assumptions to fill in the gaps

- Example:

- For a relation “people”, we keep:
 - Total number of tuples = 100,000
 - Distinct “zipcode” values that appear in it = 100
- Given a query: “zipcode = 20742”
 - We estimated the number of matching tuples as: $100,000/100 = 1000$
- What if I wanted more accurate information ?
 - Keep histograms...


Histograms

- A condensed, approximate version of the “frequency distribution”
 - Divide the range of the attribute value in “buckets”
 - For each bucket, keep the total count
 - Assume uniformity within a bucket




Histograms

- Given a query: zipcode = “ 20742”
 - Find the bucket (Number 3)
 - Say the associated count = 45000
 - Assume uniform distribution within the bucket: $45,000/200 = 225$



Histograms

- What if the ranges are typically not full ?
 - ie., only a few of the zipcodes are actually in use ?
- With each bucket, also keep the number of zipcodes that are valid
- Now the estimate would be: $45,000/80 = 562.50$
- More Information → Better estimation



Histograms

- Very widely used in practice
 - One-dimensional histograms kept on almost all columns of interest
 - ie., the columns that are commonly referenced in queries
 - Sometimes: multi-dimensional histograms also make sense
 - Less commonly used as of now
- Two common types of histograms:
 - Equi-depth
 - The attribute value range partitioned such that each bucket contains about the same number of values
 - Equi-width
 - The attribute value range partitioned in equal-sized buckets
 - others...

Estimating sizes of the results of various operations

- Guiding principle:
 - Use all the information available
 - Make uniformity and randomness assumptions otherwise
 - Many formulas, but not very complicated...
 - In most cases, your initial intuition is probably correct

Basic statistics

- Basic information stored for all relations

- $n_r = |r|$ = number of tuples in a relation r .
- b_r : number of blocks containing tuples of r .
- l_r : size of a tuple of r .
- f_r : blocking factor of r — i.e., the number of tuples of r that fit into one block.
- $V(A, r)$: number of distinct values that appear in r for attribute A ; same as the size of $\Pi_A(r)$.
- $MAX(A, r)$: the maximum value of A that appears in r
- $MIN(A, r)$
- If tuples of r are stored together physically in a file, then:

$$b_r = \left\lceil \frac{n_r}{f_r} \right\rceil$$

Selection Size Estimation

- $\sigma_{A=v}(r)$
 - $n_r / V(A, r)$: number of records that will satisfy the selection
 - equality condition on a key attribute: *size estimate* = 1
- $\sigma_{A \leq v}(r)$ (case of $\sigma_{A \geq v}(r)$ is symmetric)
 - Let c denote the estimated number of tuples satisfying the condition.
 - If $\min(A, r)$ and $\max(A, r)$ are available in catalog
 - $c = 0$ if $v < \min(A, r)$
 - $$c = n_r \cdot \frac{v - \min(A, r)}{\max(A, r) - \min(A, r)}$$
 - If histograms available, can refine above estimate
 - In absence of statistical information c is assumed to be $n_r / 2$.

Size Estimation of Complex Selections

- **selectivity**(θ_i) = the probability that a tuple in r satisfies θ_i .
 - If s_i is the number of satisfying tuples in r , then selectivity (θ_i) = s_i/n_r
- **conjunction**: $\sigma_{\theta_1 \wedge \theta_2 \wedge \dots \wedge \theta_n}(r)$. Assuming independence, estimate of tuples in the result is:

$$n_r * \frac{s_1 * s_2 * \dots * s_n}{n_r^n}$$

- **disjunction**: $\sigma_{\theta_1 \vee \theta_2 \vee \dots \vee \theta_n}(r)$. Estimated number of tuples:

$$n_r * \left(1 - \left(1 - \frac{s_1}{n_r} \right) * \left(1 - \frac{s_2}{n_r} \right) * \dots * \left(1 - \frac{s_n}{n_r} \right) \right)$$

- **negation**: $\sigma_{\neg \theta}(r)$. Estimated number of tuples: $n_r - \text{size}(\sigma_{\theta}(r))$

of tuples produced in **worst case**:

R JOIN S: $R.a = S.a$

a is primary key of S


$|R| = 10,000$; $|S| = 5000$

A. 5000

B. 10000

C. 15000

D. 50,000,000



Joins

- $R \text{ JOIN } S: R.a = S.a$
 - $|R| = 10,000; |S| = 5000$
- CASE 1: a is key for S
 - *Worst case: each tuple of R joins with exactly one tuple of S*
 - So: $|R \text{ JOIN } S| = |R| = 10,000$
- CASE 2: a is key for R
 - Similar --- so $|S| = 5,000$

Joins

- R JOIN S: $R.a = S.a$
 - $|R| = 10,000$; $|S| = 5000$
- CASE 3: a is not a key for either
 - Reason with the distributions on a
 - Say: the domain of a : $V(A, R) = 100$ (the number of distinct values a can take)
 - THEN, *assuming uniformity*
 - For each value of a
 - We have $10,000/100 = 100$ tuples of R with that value of a
 - We have $5000/100 = 50$ tuples of S with that value of a
 - All of these will join with each other, and produce $100 * 50 = 5000$
 - So total number of results in the join:
 - $5000 * 100$ (distinct values) = 500,000
 - We can improve the accuracy if we know the distributions on a better
 - Say using a histogram

Equivalence of Expressions

- Two relational expressions equivalent iff:
 - Their result is identical on all legal databases
- Equivalence rules:
 - Allow replacing one expression with another
- Examples:

1. $\sigma_{\theta_1 \wedge \theta_2}(E) = \sigma_{\theta_1}(\sigma_{\theta_2}(E))$

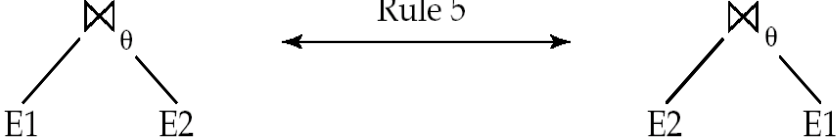
2. Selections are commutative

$$\sigma_{\theta_1}(\sigma_{\theta_2}(E)) = \sigma_{\theta_2}(\sigma_{\theta_1}(E))$$

Equivalence Rules


- Examples:

$$3. \quad \Pi_{L_1}(\Pi_{L_2}(\dots(\Pi_{L_n}(E))\dots)) = \Pi_{L_1}(E)$$

$$5. \quad E_1 \bowtie_{\theta} E_2 = E_2 \bowtie_{\theta} E_1$$


7(a). If θ only involves attributes from E_1

$$\sigma_{\theta\theta}(E_1 \bowtie_{\theta} E_2) = (\sigma_{\theta\theta}(E_1)) \bowtie_{\theta} E_2$$



- And so on...

- Many rules of this type
- See textbook for more examples

Example

- Find the names of all customers with an account at a Brooklyn branch whose account balance is over \$1000.


$$\Pi_{customer_name}(\sigma_{branch_city = \text{"Brooklyn"} \wedge balance > 1000} (branch \bowtie (account \bowtie depositor)))$$

- Apply the rules one by one

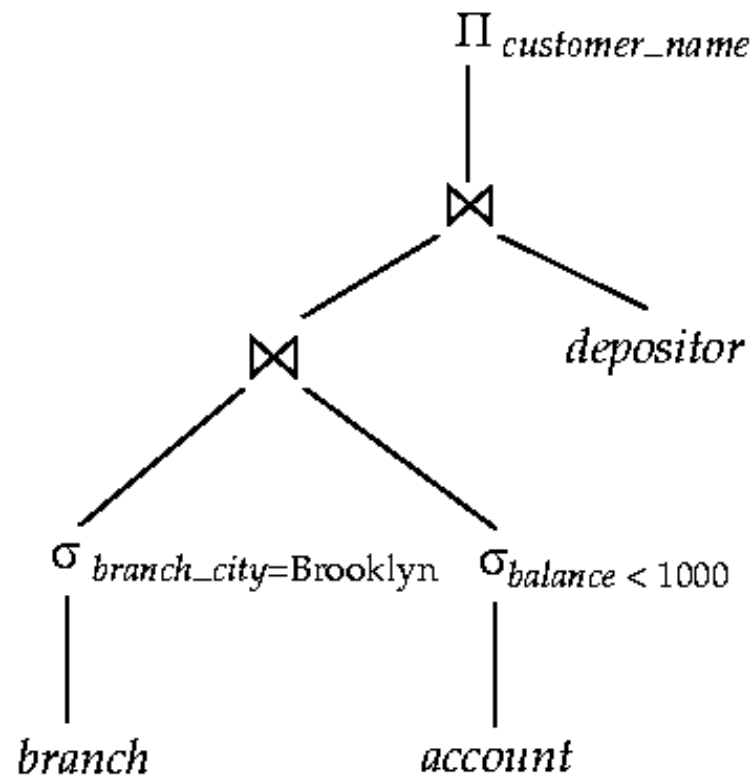
$$\Pi_{customer_name}((\sigma_{branch_city = \text{"Brooklyn"} \wedge balance > 1000} (branch \bowtie account)) \bowtie depositor)$$

$$\Pi_{customer_name}(((\sigma_{branch_city = \text{"Brooklyn"}}(branch)) \bowtie (\sigma_{balance > 1000}(account))) \bowtie depositor)$$

Example



(a) Initial expression tree



(b) Tree after multiple transformations

Equivalence of Expressions

- The rules give us a way to enumerate all equivalent expressions
 - Note that the expressions don't contain physical access methods, join methods etc...
- Simple Algorithm:
 - Start with the original expression
 - Apply all possible applicable rules to get a new set of expressions
 - Repeat with this new set of expressions
 - Till no new expressions are generated

Equivalence of Expressions

- Works, but is not feasible
- Consider a simple case:
 - $R1 \bowtie (R2 \bowtie (R3 \bowtie (\dots \bowtie Rn)))\dots$
- Just join commutativity and associativity can give us up to $n! * 2^n$ plans to consider
 - Typically enumeration combined with the search process

Evaluation Plans

- We still need to choose the join methods etc..
 - Option 1: Choose for each operation separately
 - Usually okay, but sometimes the operators interact
 - Consider joining three relations on the same attribute:
 - $R1 \bowtie_a (R2 \bowtie_a R3)$
 - Best option for R2 join R3 might be hash-join
 - But if $R1$ is sorted on a , then *sort-merge join* is preferable
 - Because it produces the result in sorted order by a
- Also, pipelining or materialization
- Such issues typically arise when doing the optimization

Query Optimization

- Integral component of query processing
- One of the most complex pieces of code in a database system
- Active area of research
 - E.g. JSON Query Optimization
 - What if you don't know anything about the statistics
 - Better statistics
 - How to prune search space
 - How good is good enough?