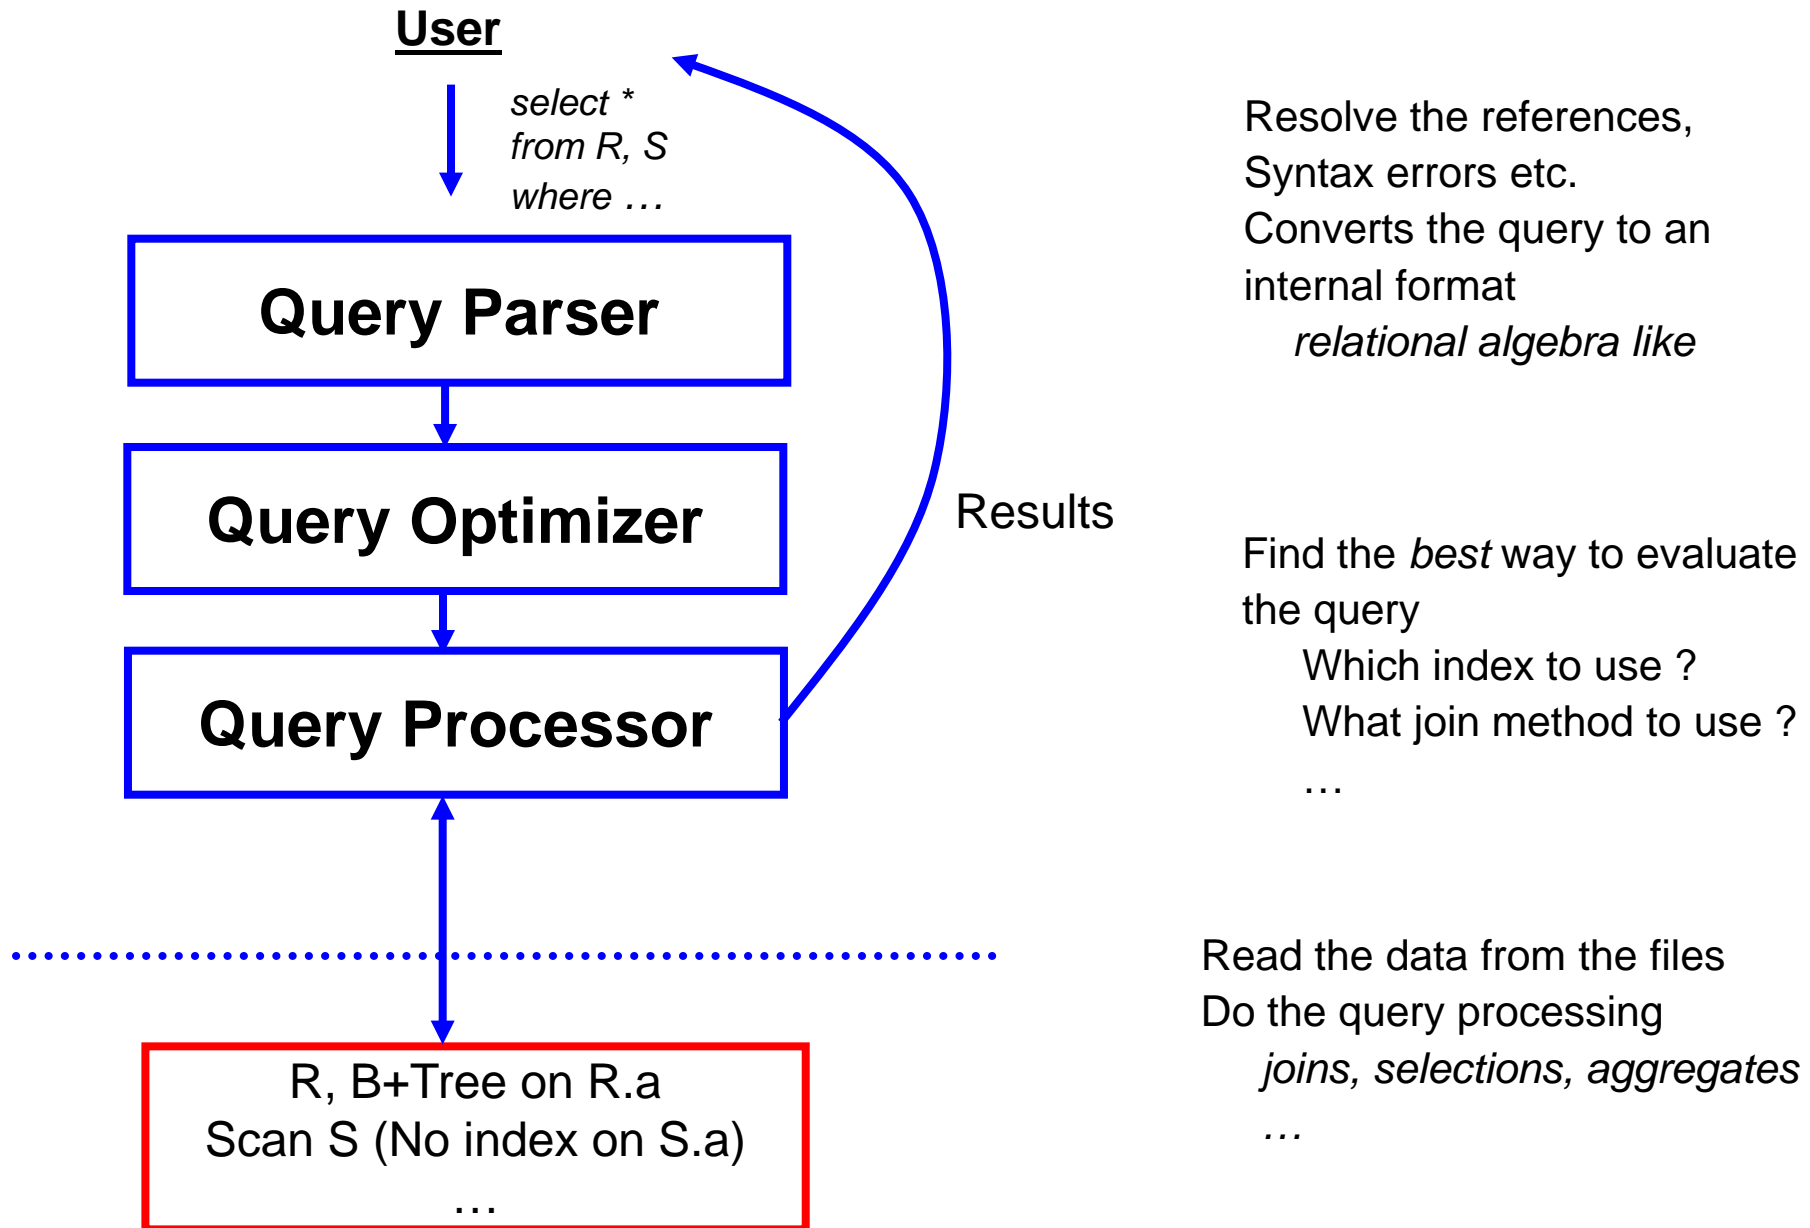


Query Processing

Overview

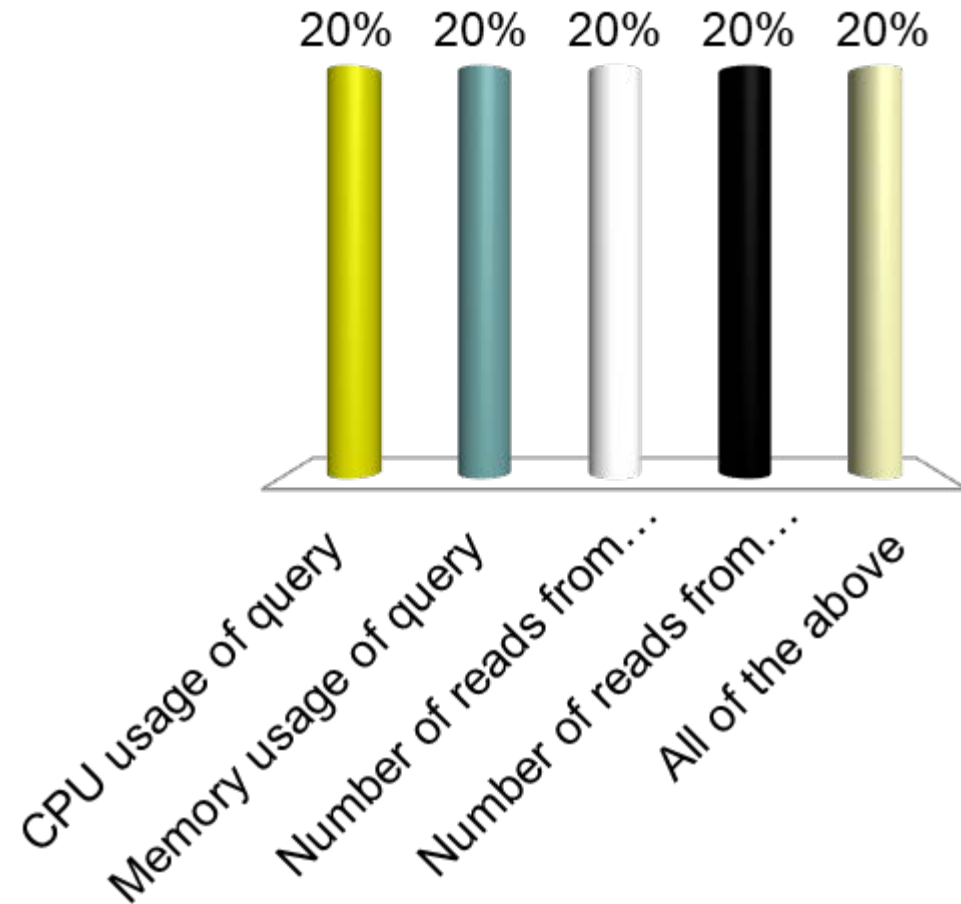


Query planning

- Performed primarily by the optimizer
- SQL query just says “what” to get, we need to figure out “how”
- Basic process:
 - Enumerate different options
 - Assign costs to different options
 - Choose lowest cost
- Cost is not the same thing as response time


In theory what is relevant in calculating cost?

- A. CPU usage of query
- B. Memory usage of query
- C. Number of reads from memory
- D. Number of reads from disk
- E. All of the above



If we can only take into account one factor, which one should we choose?

- A. CPU usage of query
- B. Memory usage of query
- C. Number of reads from memory
- D. Number of reads from disk
- E. Number of disk seeks



“Cost”

- Complicated to compute
- We will focus on disk:
 - Number of I/Os ?
 - Not sufficient
 - Number of seeks matters a lot... why ?
 - t_T – time to transfer one block
 - t_S – time for one seek
 - Cost for b block transfers plus S seeks
$$b * t_T + S * t_S$$
 - Measured in *seconds*


Selection Operation

- `SELECT * FROM person WHERE SSN = "123"`
- Option 1: Sequential Scan
 - Read the relation start to end and look for "123"
 - Can always be used (not true for the other options)
 - Cost ?
 - *Let b_r = Number of relation blocks*
 - Then:
 - 1 seek and b_r block transfers
 - So:
 - $t_S + b_r * t_T$ sec

How does result change if predicate on candidate key?

`SELECT * FROM person WHERE SSN = "123"`

- A. There is always an index on a candidate key, and we should use that instead.
- B. Since the data is sorted, we can use binary search.
- C. We know for sure that the predicate will fail, so the cost is 0.
- D. Once we hit the first result, we can stop. So we can assume cost is approximately half of the amount from the previous slide.

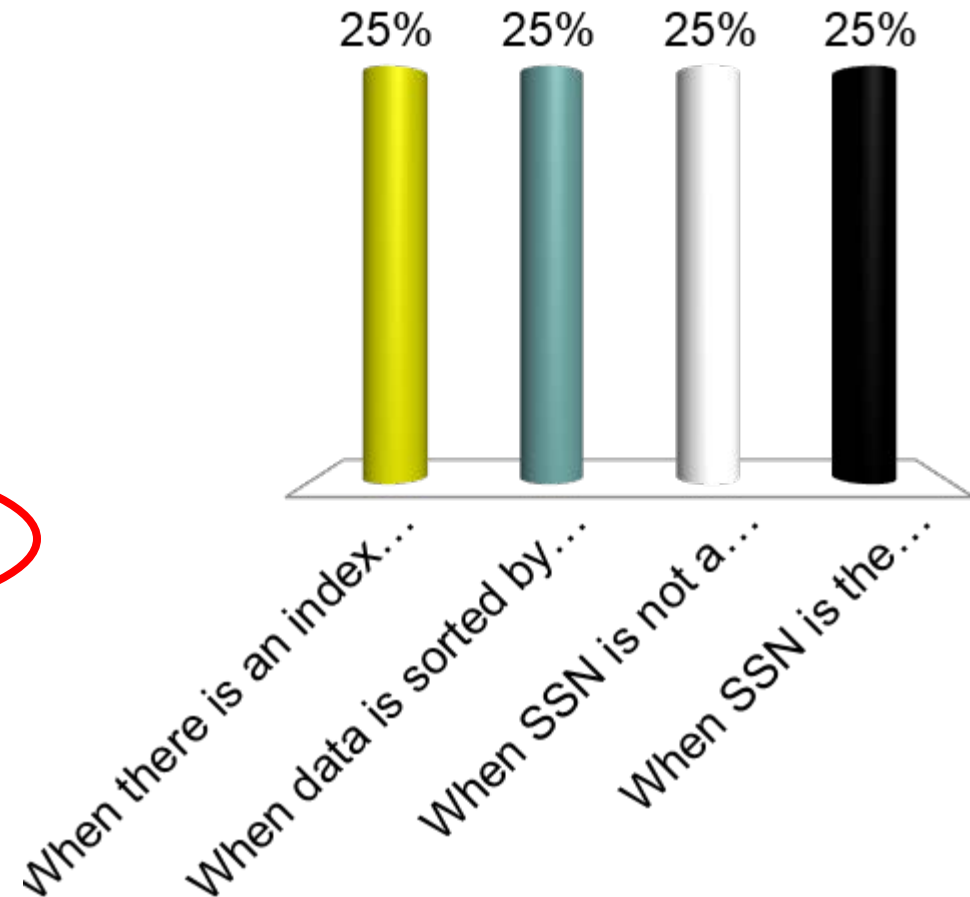


Selection Operation

- `SELECT * FROM person WHERE SSN = "123"`
- Option 2 : Binary Search:
 - Pre-condition:
 - *The relation is sorted on SSN*
 - *Selection condition is an equality*
 - E.g. can't apply to "*Name like '%424%'*"
 - Do binary search
 - Cost of finding the *first* tuple that matches
 - $\lceil \log_2(b_r) \rceil * (t_T + t_S)$
 - *All I/Os are random, so need a seek for all*
 - The last few are short hops, but we ignore such small effects

When is $\log_2(b_r) * (t_T + t_S)$ (from previous slide) too low of an estimate?

- A. When there is an index on SSN
- B. When data is sorted by SSN
- C. When SSN is not a candidate key
- D. When SSN is the primary key



Selection Operation

- `SELECT * FROM person WHERE SSN = "123"`
- Option 3 : Use Index
 - Pre-condition:
 - *An appropriate index must exist*
 - Use the index
 - Find the first leaf page that contains the search key
 - Retrieve all the tuples that match by following the pointers
 - If primary index, the relation is sorted by the search key
 - Go to the relation and read blocks sequentially
 - If secondary index, must follow all pointers using the index

Selection w/ B+-Tree Indexes

<div>why?</div>	cost of finding the first leaf	cost of retrieving the tuples
primary index, candidate key , equality	$h_i * (t_T + t_S)$	$1 * (t_T + t_S)$
primary index, not a key, equality	$h_i * (t_T + t_S)$	$1 * (t_T + t_S) + (b - 1) * t_T$ <i>Note: primary == sorted</i> <i>b = number of pages that contain the matches</i>
secondary index, candidate key , equality	$h_i * (t_T + t_S)$	$1 * (t_T + t_S)$
secondary index, not a key, equality	$h_i * (t_T + t_S)$	$n * (t_T + t_S)$ <i>n = number of records that match</i> This can be bad

h_i = height of the index

Selection Operation

- Selections involving ranges
 - *select * from accounts where balance > 100000*
 - *select * from matches where matchdate between '10/20/06' and '10/30/06'*
 - Option 1: Sequential scan
 - Option 2: Using an appropriate index
 - Can't use hash indexes for this purpose

Selection Operation

- Complex selections
 - Conjunctive: *select * from accounts where balance > 100000 and SSN = "123"*
 - Disjunctive: *select * from accounts where balance > 100000 or SSN = "123"*
 - Option 1: Sequential scan
 - Option 2 (*Conjunctive only*): Using an appropriate index on one of the conditions
 - E.g. Use SSN index to evaluate SSN = "123". Apply the second condition to the tuples that match
 - Or do the other way around (if index on balance exists)
 - Which is better ?
 - Option 3 (*Conjunctive only*) : Choose a multi-key index
 - Not commonly available

Selection Operation

- Complex selections
 - Conjunctive: *select * from accounts where balance > 100000 and SSN = "123"*
 - Disjunctive: *select * from accounts where balance > 100000 or SSN = "123"*
 - **Option 4:** Conjunction or disjunction of *record identifiers*
 - Use indexes to find all RIDs that match each of the conditions
 - Do an *intersection* (for conjunction) or a *union* (for disjunction)
 - Sort the records and fetch them in one shot
 - Called "Index-ANDing" or "Index-ORing"
 - Heavily used in commercial systems